

Basketball Prospect Predictor

By Matthew, Kincaid, Ethan

Problem Statement

- **Question:** Can we predict NBA career success based on NCAA performance and qualitative analysis?
- Two signal sources
 - NCAA statistics - per-game box score numbers (points, rebounds, shooting efficiency, etc...)
 - Scouting reports - written text evaluations of each prospect's skills, athleticism, and projection
- Why a multimodal investigation?
 - Statistics miss character, athleticism and things that may not reflect on the scoreboard
 - Scouting reports are subjective & qualitative and it makes it hard to compare across scouts
 - Using them together should cancel their blind spots
 - Goal: train classifier/regression models on NCAA stats, along with a NLP approach to analyze scouting reports, and see better performance through a multi-modal approach than through any individual models

Data Collection

- Sources:
 - NCAA stats
 - Scraped from NCAA website
 - Backfill using sports dataverse
 - NBA stats
 - NBA API for basic metrics
 - [Basketball Reference](#) for VORP
 - [databallr](#) for DARKO
 - Scouting Reports from nbadraft.net

What Does the Data Look Like?

- As stated we have three sources of data, NCAA college stats, NBA career stats, and scouting reports
 - NCAA College stats
 - 787 players, draft classes 2009-2023
 - Stats: PPG, RPG, APG, FG%, 3P% FT% BPG, SPG, MPG, OReb, DReb, TO, Games, Height, Position, Class
 - NBA Career Stats (Best Season of First Three Years in NBA)
 - 787 matched players (same players from NCAA set), draft classes 2009-2023
 - Stats: PTS, REB, AST, STL, BLK, FG%, 3P%, FT%, +/-, MIN, GP
 - Scouting Reports
 - 1795 players, Draft classes 2009-2026
 - Reports from scouts that detail the strengths, weaknesses, NBA comparison, and a full narrative report about the player
- All three sources are joined on player name & draft year to form the full multi-modal training dataset

Regression

- Model Selection

- NCAA stats produce a continuous outcome (role z-score) so regression seemed like the natural fit
- We chose three variants:
 - **Lasso**
 - We have about 18 stats and they are highly correlated, the L1 penalty automatically zeros out redundant features (built-in feature selection)
 - **Ridge**
 - Softer comparison that shrinks weights rather than zeroing them (better when features contribute a little)
 - **XGBoost**
 - Captures non-linear interactions linear models miss (ex. High efficiency & high assist rate together predicting succeeds better than either alone)


- Model Implementation

- **Input**
 - NCAA stats → Imputer → StandardScaler → $x \in \mathbb{R}^d$
 - Position → OneHotEncoder([Guard, Forward, Center])
 - Class yr → OrdinalEncoder → StandardScaler
 - Stats: pts_pg, reb_pg, ast_pg, blk_pg, stl_pg, fg_pct, ft_pct, fg3_share, ft_rate, pts_per_fga, height_dev, team_diffculty
- **Target — nba_role_zscore**
 - $z = \sum w_i \cdot \text{clip}(z_i, \pm 2.5) / \sum 1[s \text{ available}]$
 - Weights: MIN=0.30, pts=0.25, reb=0.20,
 - ast=0.15, stl=0.05, blk=0.05
 - No NBA data → floor of -3.0
- **Lasso**
 - $L = \|y - Xw\|_2^2 + \alpha \|w\|_1$
 - α : 100 log-spaced values [1e-4, 100], 5-fold CV
 - Retained only 3 features: RPG, BPKG, FG%
- **Ridge**
 - $L = \|y - Xw\|_2^2 + \alpha \|w\|_2^2$
 - Same α search — keeps all features, shrinks weights
- **XGBoost**
 - $\hat{y} = \sum f_i(x)$
 - GridSearchCV: depth [2,3], lr [0.05,0.1],
 - n_estimators [100,200], L1/L2 regularization
- **Output**
 - $\hat{y} \in \mathbb{R}$ (predicted nba_role_zscore)
 - > 1.5 = Star | $0.5-1.5$ = Starter
 - $-0.5-0.5$ = Bench | < -0.5 = Bust

Regression

● Model Performance

Model	R^2	RMSE	MAE
Lasso	0.0479	0.9811	0.7608
Ridge	-1.562	1.609	0.8405
XGBoost	0.2934	0.8452	0.7035



● Limitations:

- NBA Outcomes depend on team fit, injuries, and player development which aren't in the NCAA data
- Raw counting stats are position cofounded (center vs guard RPG mean very different things)
- Noisy continuous target (binned classification is more tractable for this dataset)

● Future Directions

- Position relative input normalization (z-score each stat within position group)
- Draft class normalization to control for era effects
- Feed continuous score into multimodal meta model alongside classification probability outputs

Classification


- Why classification?
- Tiered by distributions of key stats

- Bust: $z < -0.5$
- Bench: $-0.5 \leq z < 0.5$
- Starter $0.5 \leq z < 1.5$
- Star $z \geq 1.5$

- Logistic Regression (L1/L2)

- Extension of linear models to classification

- Softmax →



- XGBoost

- Tier boundaries not linearly separable
- Tree structure refines decision making process by capturing combinations of features

- Input

- NCAA stats → median → StandardScaler → $x \in \mathbb{R}^d$
- Position → OneHotEncoder([G, F, C]) (optional)
- Class Yr → OrdinalEncoder → Standard Scaler

Logistic Regression — L1 & L2

$$\hat{p}(y = k | x) = \frac{\exp(w_k^\top x + b_k)}{\sum_j \exp(w_j^\top x + b_j)}$$

$$\mathcal{L} = - \sum_i \log \hat{p}(y_i | x_i) + \frac{1}{C} \|w\|_p$$

- $C \in \text{logspace}(-3, 2, 20)$, selected by 5-fold CV; metric: F_1 -macro
- L1: solver = SAGA L2: solver = LBFGS

XGBoost

$$\hat{p} = \text{softmax} \left(\sum_t f_t(x) \right), \quad \text{objective: multi:softprob}$$

- GridSearchCV: $d \in \{2, 3\}$, $\eta \in \{0.05, 0.1\}$, $T \in \{100, 200\}$, $s \in \{0.7, 0.9\}$

- Class imbalance — essential



- $w_i = N / (K \times N_{\square})$
- N = total, K = num classes, $N_{\square}$ = class count

- Output

- $\hat{p} = [p_{\text{bust}}, p_{\text{bench}}, p_{\text{starter}}, p_{\text{star}}] \in [0,1]$

Classification

Model	Accuracy	F1-macro	Bal.Acc
LogisticL1	0.3228	0.3154	0.3528
LogisticL2	0.3228	0.3160	0.3547
XGBoost	0.4241	0.4365	0.4523



Limitations

- Personality, work ethic, injuries, etc
- Arbitrary boundary decisions and player classification (what makes a player 'good')
 - Impact vs presence
- Star class underrepresented
- Ordinal structure ignored
 - Adjacent errors should be penalized less than distant ones

Future Directions

- Ordinal regression so we can respect rank ordering of the tiers
- Probability Calibration
 - Add isotonic regression so probability outputs are reliable for meta model input
- Independent study - what makes an NBA player 'good'

Text – Version 1

● Model Selection

- Uses DistilBERT on scouting reports
 - Short but context-heavy text
 - Captures nuanced language beyond box score
- Report → compact deep embedding
- Can handle regression + classification
 - Z score, VORP, DARKO
 - Prospect tier
 - Derived from Z score

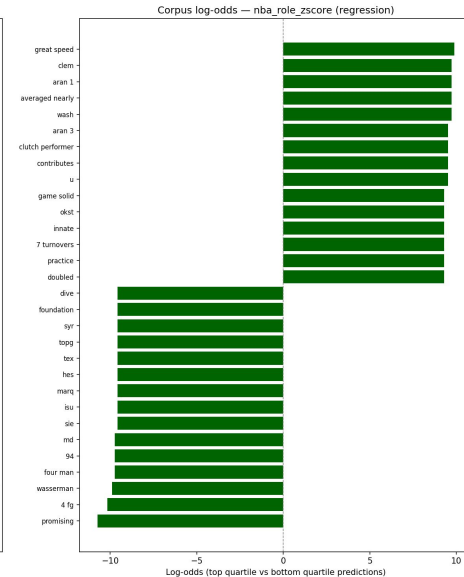

● Model Implementation

- Tokenized and encoded with DistilBERT
 - Batch size of 32
- Pooled text representation into:
 - MLP block: Linear(768 → 64) → ReLU → Dropout
 - Shared hidden layer: Linear(64 → 32) → ReLU → Dropout
- Task-specific output head:
 - Regression output: one continuous value per player
 - Linear(32 → 1)
 - Classification output: 4 class logits (converted to probabilities)
 - Linear (32 → 4)

Text – Version 1

- Interpretation Approach


- Gut check... what does a report for a good NBA player look like
- Use chosen phrases to test how specific scouting language shifts model output
- Compare language patterns in high-prediction vs low-prediction reports
- Analyze tone across strengths, weaknesses, and outlook sections against predictions



Text – Version 1

- Model Performance

- Results were not good
 - Regression: MAE $\approx$ 0.92
 - Classification: Accuracy $\approx$ 0.29, Macro-F1 $\approx$ 0.17
- Model captures *some* signal,
 - Not strong enough as a standalone predictor



- Limitations + Future Direction

- Scouting language is noisy and inconsistent across sources
 - A lot of irrelevant information
- Text coverage and detail vary by player, which weakens supervision quality
- Use interpretability outputs to identify which phrasing patterns are actually predictive
- Primary role going forward: complementary feature source, not standalone model


Text – Version 2

● Model Selection

- The idea: maybe simpler model will perform better
 - Less prone to noise distraction?
- Shallow model focused on speed and transparency
- Uses sentiment and scouting-term signals

● Model Implementation


- Extract VADER sentiment
- Count set of good/bad words/phrases
- Train Ridge regression model
- Convert into tier probabilities



Text – Version 2

- Model Performance

- Results were not good
 - $MAE \approx 0.90$, $R^2 \approx 0.00$
- Again, there is *some* signal
 - But not very strong



- Limitations + Future Direction

- Handcrafted language features can miss context and subtle phrasing
- Feature quality depends on lexicon coverage and report consistency
- Goal is not text-only dominance, but better complementary signal for final stacking

```
[[log_reg: p_bust, p_bench...],  
[xgb_reg: p_bust, p_bench...],  
[log_cl: p_bust, p_bench...],  
[xgb_cl: p_bust, p_bench...],  
[nlp: p_bust, p_bench...]]
```

Multi-Modal Meta-Model

● Why MMMM

- Classification optimizes for boundaries
- Regression captures **where** in bound a player is expected to fall
- Text (in theory) captures non-statistical factors (attitude, work ethic, etc)

● Unified Probability Interface

- Base models output in the same format
- Regression z-scored converted to probabilities using Gaussian CDF
- Enables modular base model swapping via config

● Stacking Architecture

- Base model $p(k)$ concatenated into a meta-feature vector
- Logistic regression (meta-model) learns which base models to trust for which tiers
- Weights = reliability of each model

● Out-of-Fold (OOF) Training


- Base model trained on $k-1$ folds and predicts held out fold
- MMMM trained only on OOF predictions
- Prevents MMMM from seeing same training data 4 times

Meta-Model Results

Confusion Matrix

Rows = actual tier | Columns = predicted tier

Cell: count and % of that actual tier predicted as each column tier




Multimodal Model Comparison

(blue row = final stacker output; other rows = individual base model performance)



Model	Acc	Within 1	Dist Wt Acc	Ord MAE	Exp MAE	QWK	F1 Macro
Stacker (final)	0.4375	0.8413	0.7532	0.7404	0.7191	0.4902	0.4312
CLF - L2 Logistic	0.3558	0.7548	0.6987	0.9038	0.7476	0.1818	0.3442
CLF - XGBoost	0.4135	0.7885	0.7276	0.8173	0.7453	0.1996	0.3285
REG - Ridge	0.3846	0.9279	0.7708	0.6875	0.7141	0.3801	0.2827
REG - XGBoost	0.3846	0.9519	0.7788	0.6635	0.6728	0.4544	0.3047
text_scouting	0.3798	0.9087	0.7628	0.7115	0.8425	0.1807	0.2166
text_scouting_lexico	0.3990	0.8654	0.7516	0.7452	0.8399	0.2203	0.2631

Prediction Error Distribution

How many tiers off was the model? (lower error = better)



Meta-Model Results



Future Direction

1. Improving data processing approach
2. Getting combine data as a proxy for physicality & health
3. Implementing proprietary classification/regression models to respect ordinality
4. Further exploring NLP and different scouting sources
5. Hyperparameter optimization
6. Conflict between predicting a continuous output and placing into buckets

Predicting the 2026 NBA Draft

